Trellis: Full Review of BRD, SRS and Architecture

Oct 23, 2026 · @Hussain

1. Overview and how to read this review

The three Trellis documents are well structured and unusually honest about uncertainty, but they freeze about 70 requirements before the Instagram, competitor-data and free-AI feasibility questions are answered. This review gives feedback on every section, requirement, table row, figure, decision and gate in the BRD v1.0, SRS v1.0 and Architecture v0.1, whether the feedback is positive, neutral or negative.

Each item is rated with one of four tags:

- **Strong**: keep as is, or very close to it.
- **Adequate**: sound intent, small fixes needed.
- **Weak**: needs rework before it can be built or tested.
- **Gap**: something missing or unspecified that the item implies.

API and vendor facts come from general knowledge of Meta's platform and free AI tiers, not a fresh check of current terms. Treat them as claims for the feasibility spikes to confirm.

Document	Items reviewed	Strong	Adequate	Weak / Gap
BRD v1.0	sections, 6 objectives, 14 BRs, 8 rules, 5 criteria, 4 changes, 5 assumptions	about 20	about 18	about 12
SRS v1.0	8 AI, 9 CONT, 8 DATA, 4 COMP, 5 DISC, 1 NEWS, 6 PLAN, 3 NOTIF, 8 AUTH, 11 NFR, 10 tests, 8 conflicts	about 28	about 40	about 28
Architecture v0.1	6 principles, 9 modules, 8 data domains, 3 flows, 8 ADRs, 6 risks, 7 gates, 4 spikes, 3 figures	about 25	about 25	about 15

The counts are approximate and only there to show where the weight sits: the BRD is the cleanest, the SRS is broadest but least testable, and the architecture is sound in intent but thin on models and mechanisms.

2. BRD sections 1 to 3: purpose, objectives, stakeholders, scope

The opening sections are clear and honest about the document's status, but the objectives are not measurable and several scope items are really implementation constraints.

2.1 Purpose and authority (BRD §1)

Item	What it says	Rating	Feedback
Product definition	AI-assisted Instagram growth, analysis and content-planning web app for one skincare/K-beauty creator, with possible commercialization	Strong	A one-sentence product definition that names the niche and the audience. The commercialization clause is useful but quietly drives later conflicts (free-tier AI terms, Meta App Review).
Product Owner is the first user	PO is the user; project lead is Scrum Master and technical decision maker	Adequate	Clear roles, but "Scrum Master" for a solo build adds ceremony without value. No sprint or delivery cadence is defined anywhere, so the title has nothing to govern.
"Captures needs, not an implementation plan"	BRD scope boundary	Adequate	Good principle, then contradicted by the SRS fixing the stack and the BRD itself excluding hosting choices. See cross-document findings.
Source baseline	PO Feature Requirements v1.1 plus later decisions; later decisions logged under change control	Strong	Excellent practice: says where requirements came from and refuses to disguise later changes as original. The PO v1.1 document is not attached, so reviewers cannot verify the diff.

2.2 Business objectives (BRD §2)

Objective	Rating	Feedback
Reduce time and effort to interpret performance and plan content	Weak	The right goal but unmeasured. No baseline (how long it takes now) and no target. Needs a simple measure, for example minutes from opening the app to a decided weekly plan.
Evidence-based, personalized guidance instead of generic or overly literal AI	Adequate	Captures the real product risk. "Overly literal" is vague; say what a literal answer looks like and how it would be caught.
Natural, authentic writing in the creator's voice with editorial control	Adequate	Strong intent. "Natural" and "authentic" need an operational test (blind comparison against the creator's own captions is one option).
Practical planning and discovery on the iPad	Strong	A clear, testable device target, though no iPad model or Safari version is given.
Preserve trust through data limitations, explicit approval and account security	Strong	The best objective in the set. It is the foundation of most later design decisions.
Path toward multiple independent customers without hardcoding the initial account	Adequate	Reasonable and cheap if done early. It deserves a stated boundary: model multiple users, but don't build billing, onboarding or admin tooling.

2.3 Stakeholders and scope (BRD §3)

Item	Rating	Feedback
Stakeholder table (PO, Scrum Master/tech lead, future end user)	Weak	Two real people and one hypothetical. The table adds little; user journeys and personas would add more. The future end user is described only by ownership, with no needs.
In scope: responsive browser app, iPad first, equivalent core functions elsewhere	Adequate	"Equivalent core functions" on three form factors is a large cost and is never priced. See BR-14.

Item	Rating	Feedback
In scope: Dashboard auditing, analytics, historical comparisons, opportunities, in-Dashboard planner	Adequate	Clear list, but "auditing" is never defined.
In scope: AI Chat for contextual analysis and image-led caption generation	Adequate	Two different products (analyst and writer) share one line; they have different quality risks.
In scope: competitor identification and comparative analysis	Weak	Dependent on data access that is unverified. The scope line should be conditional on Spike B.
In scope: combined Hot Topics, Niche Ideas and swipe discovery	Weak	No identified data source for "Hot Topics". The scope is stated without a feasible input.
In scope: official Instagram algorithm and Meta Edits updates	Adequate	Valuable and small, but the sources are not machine-friendly. Needs a named feed or source list.
In scope: sign-in, explicit connection confirmation, secure handling, consent-based retention	Strong	Right items and appropriate for a trust-centred product.
In scope: user-initiated one-click publishing and in-app reminders	Adequate	A sensible reduction from automatic scheduling. Reminders only in-app is the weak half.
Excluded: unattended publishing, retries, native apps, multi-account, paid AI, "My Brand", forecast guarantees, hosting choices, enterprise infrastructure	Adequate	A good exclusions list, which many specs lack. "My Brand" is never defined, and "hosting choices" is excluded even though the free-only budget makes hosting the largest hidden constraint.

3. BRD section 4: business requirements BR-01 to BR-14

Most BRs are feature descriptions rather than business needs, but they are concise and map cleanly onto SRS groups. None has a priority, an owner or a measure.

ID	What it says	Rating	Feedback
BR-01	Audit profile, bio, posts, audience and performance with actionable opportunities	Adequate	Core value of the product. "Audit" and "actionable" are undefined; the data available (especially audience data) depends on Meta permissions.
BR-02	Compare with historical periods; investigate follower changes without speculation stated as fact	Strong	The "no speculation as fact" clause is exactly right. It needs a data-collection requirement behind it, because history only exists if snapshots are taken.
BR-03	Assistant understands natural language, selects evidence, gives grounded answers	Strong	A clear statement of the main product bet. Needs an evaluation method to know whether it is met.
BR-04	Image-informed hooks, captions, CTAs and hashtags in an editable fill-in-the-blanks workflow	Adequate	Good, concrete workflow. Depends on free-tier vision quality, which is unverified.
BR-05	Text reflects past captions and experience; competitor material informs strategy, not voice	Strong	A clear ethical and quality line. "Reflects experience" slightly conflicts with the later rule not to invent experience; reconcile the wording.
BR-06	Hard prohibited-language list; AI can suggest changes but cannot activate them	Strong	Excellent control design. Needs precise matching rules (case, inflection, hashtags) at SRS level.
BR-07	Organize drafts, daily content choices and formats in a planner inside Dashboard	Adequate	Clear. Combines several features (drafts, daily choices, formats) that will need separate stories.
BR-08	Set posting date, get in-app reminder, explicitly press Publish each time	Strong	Good safety rule. Reminders reach the user only if the app is open, which weakens the value.
BR-09	All five latest posts must be skincare-related for a competitor to qualify	Weak	An implementation rule stated as a business need. Undefined "skincare-related", rigid threshold, pinned posts break "latest five".

ID	What it says	Rating	Feedback
BR-10	Combined Hot Topics, Niche Ideas and Swiper with evidence, rationale, format, uncertainty, save/pass, AI discussion	Weak	Appealing concept with no identified data source for topic activity. Bundles six capabilities into one requirement.
BR-11	Official algorithm and Edits news, distinguished from the assistant's inference	Strong	Valuable and small. The attribution requirement is the right design.
BR-12	Refresh preferences, manual refresh, understanding of measurement freshness	Adequate	Good transparency goal. The refresh trigger model conflicts with the need for historical snapshots.
BR-13	Connect the correct account via explicit confirmation; control disconnect and retention consent	Strong	A clear trust requirement. Retention options must be checked against Meta's platform terms.
BR-14	iPad-first usability with full functionality on desktop and mobile	Weak	"Full functionality" on three form factors is an unpriced commitment. Specify the iPad as fully supported and the others as functional but not optimized.

Overall: the set covers the product well, but it has no priority tags, no "why" for each need, and no acceptance measure. Add a column for priority and a column for the measurable outcome.

4. BRD sections 5 to 9: rules, success criteria, changes, assumptions, sign-off

The business rules are the strongest part of the BRD. The success criteria, assumptions and sign-off are where it is weakest, because they describe uncertainty without assigning anyone to resolve it.

4.1 Business rules and trust principles (BRD §5)

Rule	Rating	Feedback
Preserve history; weight older posts by relevance to the question, not by arbitrary exclusion	Adequate	Right principle. No weighting method is described, and the analytics must be built to support it.
Compare competitors only on comparable accessible measures; never invent missing measures	Strong	Core to credibility. Mirrored well in COMP-03 and the architecture's "unknown, not zero" rule.
Every recommendation states why it is relevant and the supporting observations, with no performance promise	Strong	Clear, testable and consistent with DISC-03.
Prohibited phrases are strict restrictions until the user removes them; other language stays available	Strong	Precise and respects user control. Leaves open how "strict" is enforced technically.
Planned dates and reminders are not permission to post; publishing needs explicit action per post	Strong	The best rule in the document. Repeated consistently in all three documents.
Each connection needs supported authorization and explicit identity confirmation	Strong	Prevents wrong-account mistakes. Fine as written.
Disconnect stops access; the user chooses whether permissible stored data is kept or deleted	Adequate	Good user control. "Permissible" is doing a lot of work; the legal position is not yet known.
Toasts are brief and auto-dismiss; dismissing does not erase the underlying task	Adequate	The second half is a real requirement. The first half is UI detail that belongs in the SRS or design.

4.2 Success criteria (BRD §6)

Criterion	Rating	Feedback
Ordinary questions like "Which of my recent posts has the highest views?" work without field names	Adequate	A concrete example, which is good. It is a scenario rather than a measure; needs a question set and a pass rate.

Criterion	Rating	Feedback
Captions revisable without invented experiences, robotic repetition or forbidden phrases	Adequate	Three testable failure modes in one line. Each needs its own check.
Every swipe idea communicates rationale and limitations	Strong	Directly testable by inspection.
No content is published merely because a planned date arrived	Strong	A clean negative test. It protects trust but does not measure value.
Saved drafts recoverable after session interruption; interface does not feel stalled during AI processing	Adequate	"Does not feel stalled" is subjective. Give a progress-indicator rule or a time bound.
Quantitative thresholds deferred to validation and testing	Weak	Deferring all numbers to testing means nothing can fail. Write initial targets as hypotheses now, then adjust.
Missing: any outcome measure (time saved, posts planned, creator satisfaction)	Gap	There is no way to tell, after building it, whether it helped. Even one subjective weekly rating would do.

4.3 Changes relative to PO summary v1.1 (BRD §7)

Change	Rating	Feedback
Daily audit refresh replaced by analysis on session start or request, with configurable metric refresh	Weak	Reasonable for cost, but it undermines historical snapshots. A scheduled snapshot is needed regardless of when the analysis runs.
Automatic scheduled publishing replaced by planned date, reminder and manual publish	Strong	A sound scope and risk reduction, well documented. Still needs the PO's signed acknowledgement, and it is repeated as pending in all three documents.
"Ask for relevant experience" replaced by editable fill-in-the-blanks content	Strong	Better product behaviour and avoids invented testimonials.

Change	Rating	Feedback
Trending and non-trending ideas replaced by observed evidence plus reasoned opportunities, no trend prediction	Strong	Honest scope reduction that prevents unfalsifiable claims.
Only four changes logged	Adequate	Good that they are tabulated. Check that the other differences from PO v1.1 (for example the Growy baseline and the free-only constraint) are also recorded.

4.4 Assumptions, dependencies and open decisions (BRD §8)

Item	Rating	Feedback
Data providers, access, quotas, pricing, coverage to be validated against Growy-referenced functionality	Weak	The most important risk, described as an assumption. Growy is undefined in the document and may rely on methods that are not compliant.
One account per user; personal versus professional account features need verification	Adequate	Correct caution. Publishing and insights generally need a professional account, so this should be stated as a likely requirement.
Navigation labels and three-section allocation retained from PO baseline but not invented	Weak	The structure of the app is missing from the requirements pack. A reader cannot see what the three sections are.
Model list, frequencies, SLAs, notification timing, retention and legal review remain open	Weak	An honest list but with no owners, dates or decision criteria.
PO must acknowledge the automatic-to-manual publishing change	Adequate	Correct, and should be resolved now rather than repeated.

4.5 Review and sign-off (BRD §9)

Item	Rating	Feedback
Status line: working baseline, not a representation that every refinement is approved	Weak	The hedge undermines the purpose of a baseline. Either approve it or label it a draft.
Signature block for Product Owner and technical lead	Adequate	Fine for a formal process. For a one-person project, a dated approval note per version is enough.

5. SRS sections 1 to 3: boundary, AI and Chat, content and creator voice

The AI and content requirements show good product judgement (no invented experiences, user-controlled restrictions) but mix product behaviour, engineering mechanism and policy in single lines, and most cannot be tested as written.

5.1 Scope and system boundary (SRS §1)

Item	Rating	Feedback
Derived from PO v1.1 plus later decisions; conflicts logged in §10	Strong	Clear lineage and a single place for conflicts.
"Documents expected behaviour, not component architecture"	Adequate	Good intent, then the same paragraph fixes Angular, Go and PostgreSQL. A stack choice does not belong in a requirements boundary without rationale.
AI reasoning is model-led with controlled backend tools and stored/live retrieval	Adequate	A reasonable architecture hint. It is a design decision presented as a boundary.
Free-tier only for AI and infrastructure; providers and hosting unselected	Weak	A strong constraint with major consequences (data-use terms, rate limits, sleeping hosts), stated without exploring any of them.

5.2 AI and Chat (AI-01 to AI-08)

ID	What it says	Rating	Feedback
AI-01	Interpret natural language, infer defaults, disclose consequential assumptions	Strong	The disclosure clause is the key trust behaviour. "Consequential" needs a definition or examples.
AI-02	Use recent chat turns, selective prior context and cross-chat memory only on explicit request	Adequate	Good privacy stance. Missing: how the user views, edits and deletes stored memories.
AI-03	Select metrics by question; retrieve authorized data through backend-controlled access	Strong	Correct separation of model and data access. Architecture follows it well.
AI-04	Ground numeric claims in retrieved measurements; deterministic computation; separate fact, inference, unknown	Strong	The most important AI rule. Needs a mechanism: structured answers citing evidence IDs, with numbers rendered by the server.
AI-05	Prefer complete validated answers over fast incomplete ones; clarify only when unresolvable	Adequate	Right priority, but unbounded: free-tier timeouts and rate limits need a maximum wait or an async path.
AI-06	Bounded compatible fallback on free-provider limits; never imply success after total failure	Adequate	Good honesty rule. "Bounded" and "compatible" are undefined (how many attempts, which models, same evidence format?).
AI-07	Classify output-check needs (calculations, sources, image/product consistency, prohibited language) without separate AI experiences	Weak	The last clause is unclear. The check types are well chosen, but there is no description of how each is performed.
AI-08	Support analytics, strategy, writing, official updates; no autonomous publish via Chat	Strong	A clean safety boundary, correctly mirrored in the architecture.

5.3 Content and creator voice (CONT-01 to CONT-09)

ID	What it says	Rating	Feedback
CONT-01	Interpret images and text into a complete editable post package; associate descriptions with the correct product/media	Adequate	A clear deliverable. Product-to-image association has no ground truth and no fallback when the model is unsure.
CONT-02	Placeholders or selectable suggestions when first-hand experience is unknown; no invented personal review	Strong	Excellent and rare. It is the core of the product's authenticity promise.
CONT-03	Edit, replace, remove and fill generated sections; review and save	Strong	Clear. Implies a structured editor component, which the stack does not mention.
CONT-04	Learn writing patterns from selected relevant historical captions	Weak	No data requirement: how many captions, how selected, and what happens for an account with few. "Learn" implies training; retrieval into the prompt is probably what is meant.
CONT-05	Competitor material as evidence for strategy, hooks and formats, not voice copying	Strong	Good ethical line. Hard to test, because the boundary between a hook idea and copying is fuzzy.
CONT-06	Manual add/edit/remove of strict prohibited items; AI may propose, only users activate	Strong	A clean control model, easy to enforce in code.
CONT-07	Check drafts against active prohibited expressions before delivery as ready-to-use	Adequate	Right behaviour. Missing: matching semantics (case, inflection, hashtags, languages) and what to do if the model keeps reintroducing a term.
CONT-08	Use rejections and edits as context, not as permanent blacklist entries	Strong	Subtle and good. It avoids silently hardening one edit into a rule.

ID	What it says	Rating	Feedback
CONT-09	Guidance on Reel, carousel or other formats with rationale	Adequate	Useful. Unlike DISC-03 there is no requirement that the rationale cite evidence.